

# **Advanced Database Systems**

COMP7/8116 (Fall 2024)


## Distributed Query Processing

Xiaofei Zhang

# Outline

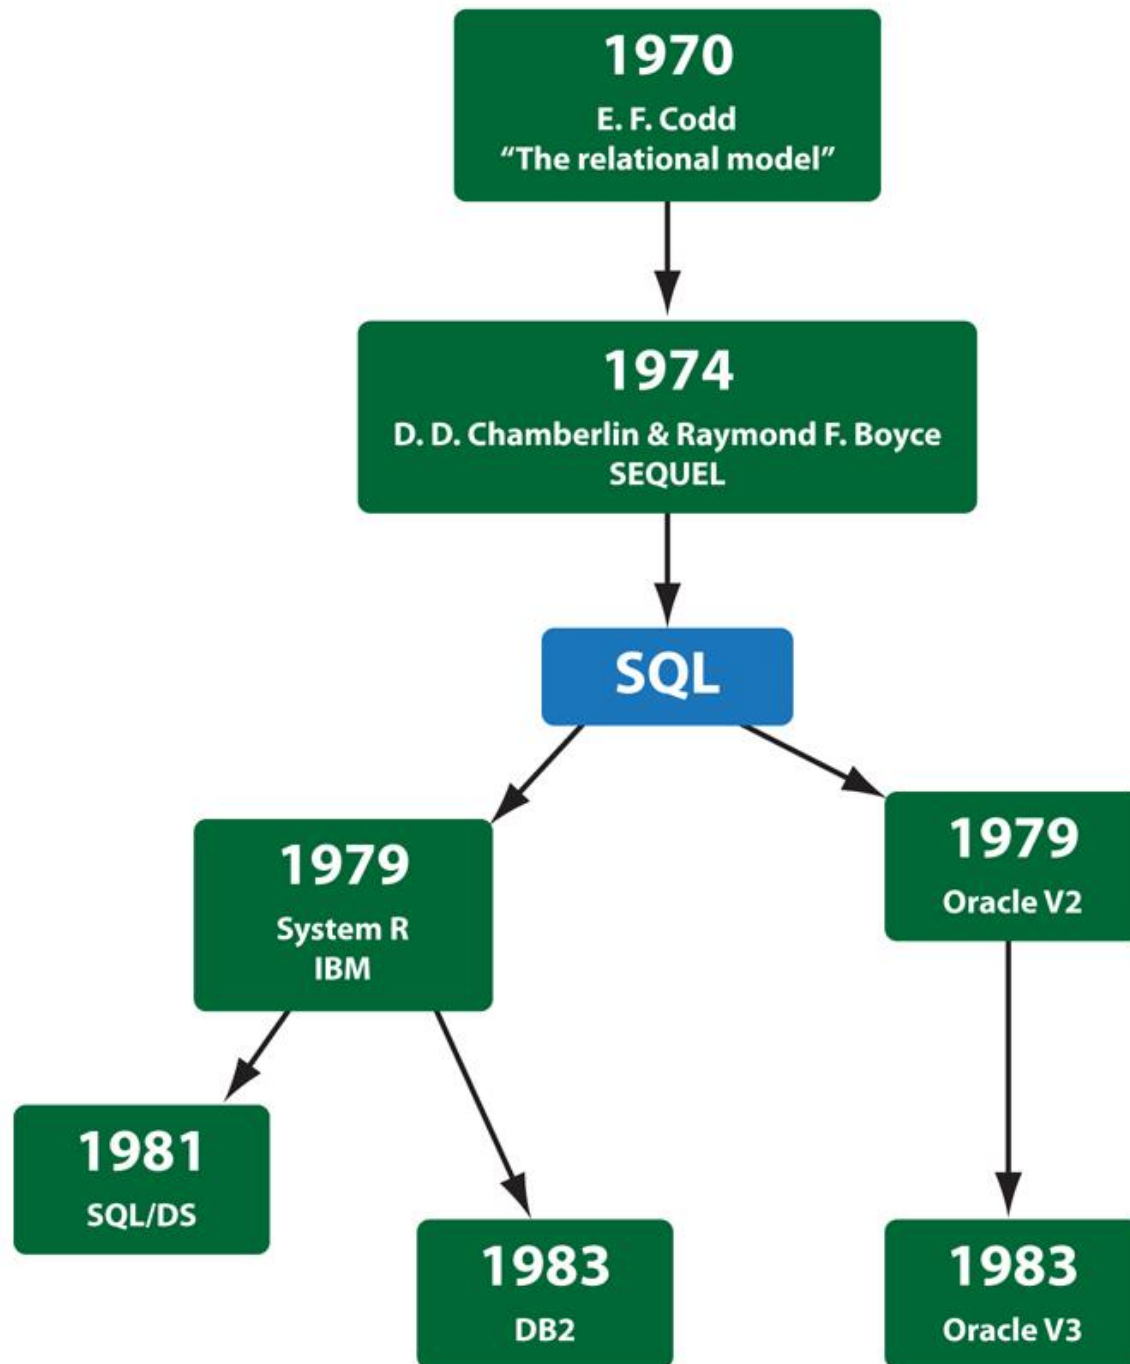
- Distributed Query Processing
  - Query Decomposition and Localization
  - Join Ordering
  - Distributed Query Optimization
  - Adaptive Query Processing

# Query Processing in a DDBMS



# Query Processing Components

- Query language
  - SQL: “intergalactic dataspeak”
- Query execution
  - The steps that one goes through in executing high-level (declarative) user queries.
- Query optimization
  - How do we determine the “best” execution plan?
- We assume a homogeneous D-DBMS



← **S.Q.L**

**IBM**  
**MariaDB**  
**MySQL**  
**PostgreSQL**

**Sequel** ↗

**Microsoft**  
**Oracle**

# Selecting Alternatives

```
SELECT  ENAME  
FROM    EMP NATURAL JOIN ASG  
WHERE    RESP = "Manager"
```

Strategy 1

$$\Pi_{ENAME}(\sigma_{RESP="Manager" \wedge EMP.ENO=ASG.ENO}(EMP \times ASG))$$

Strategy 2

$$\Pi_{ENAME}(EMP \bowtie_{ENO} (\sigma_{RESP="Manager"}(ASG)))$$

Strategy 2 avoids Cartesian product, so may be “better”

# What is the Problem?

Site 1

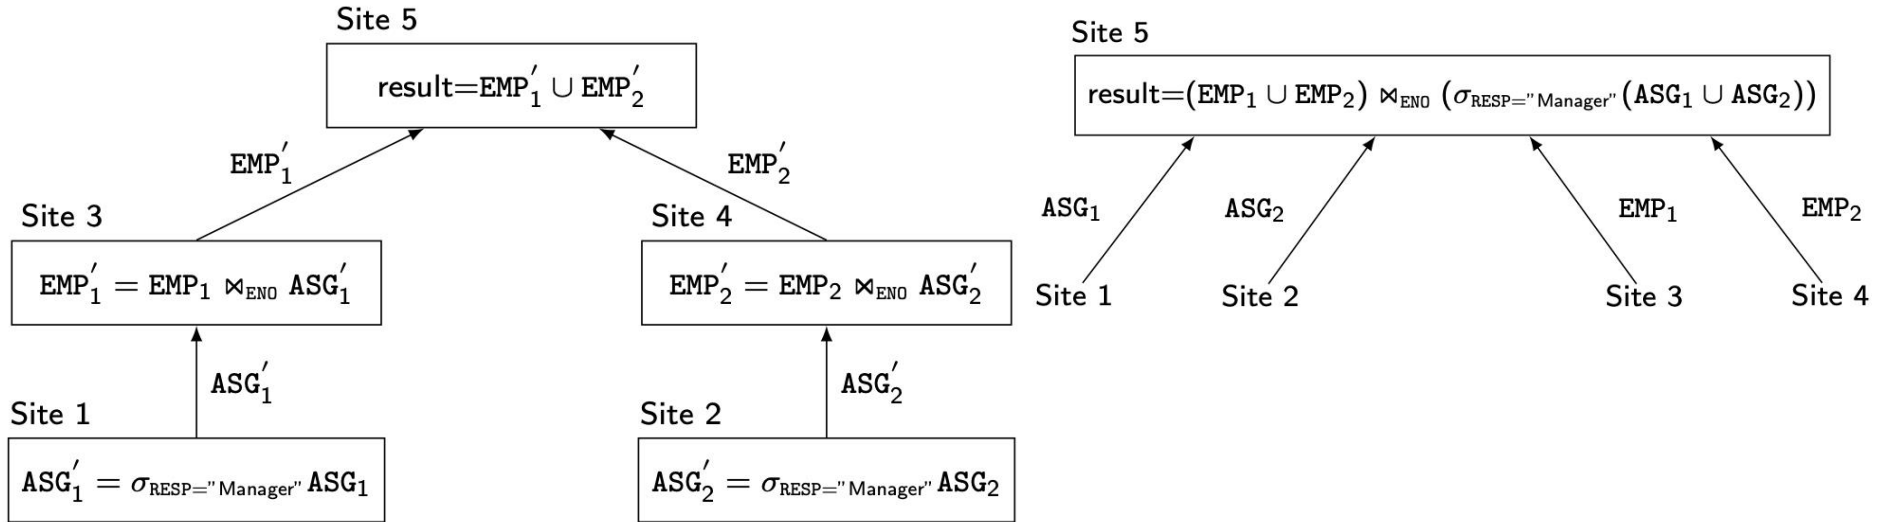
Site 2

Site 3

Site 4

Site 5

$ASG_1 = \int_{ENO \leq "E3"}(ASG)$   $ASG_2 = \int_{ENO > "E3"}(ASG)$   $EMP_1 = \int_{ENO \leq "E3"}(EMP)$   $EMP_2 = \int_{ENO > "E3"}(EMP)$  Result





# Cost of Alternatives

- Assume

- $size(EMP) = 400$ ,  $size(ASG) = 1000$
- tuple access cost = 1 unit; tuple transfer cost = 10 units
- There are 20 managers (evenly distributed in ASG)

- Strategy 1

- produce ASG':  $(10+10) * \text{tuple access cost}$  20
- transfer ASG' to the sites of EMP:  $(10+10) * \text{tuple transfer cost}$  200
- produce EMP':  $(10+10) * \text{tuple access cost} * 2$  40
- transfer EMP' to result site:  $(10+10) * \text{tuple transfer cost}$  200

Total Cost

460

- Strategy 2

- transfer EMP to site 5:  $400 * \text{tuple transfer cost}$  4,000
- transfer ASG to site 5:  $1000 * \text{tuple transfer cost}$  10,000
- produce ASG':  $1000 * \text{tuple access cost}$  1,000
- join EMP and ASG':  $400 * 20 * \text{tuple access cost}$  8,000

Total Cost

23,000

# Query Optimization Objectives

- Minimize a cost function
  - I/O cost + CPU cost + communication cost
  - These might have different weights in different distributed environments
- Wide area networks
  - Communication cost may dominate or vary much
    - Bandwidth
    - Speed
    - Protocol overhead
- Local area networks
  - Communication cost not that dominant, so total cost function should be considered
- Can also maximize throughput

# Complexity of Relational Operations

- Assume
  - Relations of cardinality  $n$
  - Sequential scan

Operation	Complexity
Select Project (without duplicate elimination)	$O(n)$
Project (with duplicate elimination) Group	$O(n * \log n)$
Join Semi-join Division Set Operators	$O(n * \log n)$
Cartesian Product	$O(n^2)$

# Types Of Optimizers

- Exhaustive search
  - Cost-based
  - Optimal
  - Combinatorial complexity in the number of relations
- Heuristics
  - Not optimal
  - Regroup common sub-expressions
  - Perform selection, projection first
  - Replace a join by a series of semi-joins
  - Reorder operations to reduce intermediate relation size
  - Optimize individual operations

# Optimization Granularity

- Single query at a time
  - Cannot use common intermediate results
- Multiple queries at a time
  - Efficient if many similar queries
  - Decision space is much larger

# Optimization Timing

- Static
  - Compilation → optimize prior to the execution
  - Difficult to estimate the size of the intermediate results®error propagation
  - Can amortize over many executions
- Dynamic
  - Run time optimization
  - Exact information on the intermediate relation sizes
  - Have to re-optimize for multiple executions
- Hybrid
  - Compile using a static algorithm
  - If the error in estimate sizes > threshold, re-optimize at run time

# Statistics

- Relation
  - Cardinality
  - Size of a tuple
  - Fraction of tuples participating in a join with another relation
- Attribute
  - Cardinality of domain
  - Actual number of distinct values
- Simplifying assumptions
  - Independence between different attribute values
  - Uniform distribution of attribute values within their domain

# Optimization Decision Sites

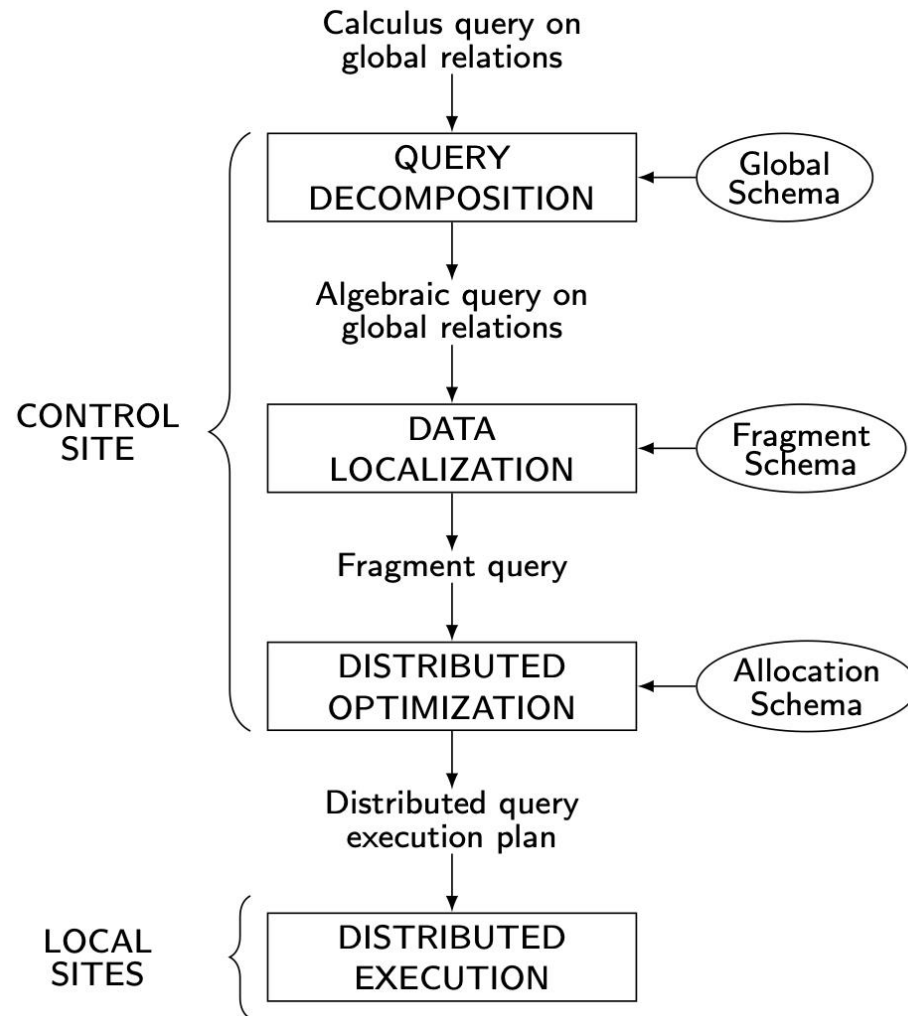
- Centralized
  - Single site determines the “best” schedule
  - Simple
  - Need knowledge about the entire distributed database
- Distributed
  - Cooperation among sites to determine the schedule
  - Need only local information
  - Cost of cooperation
- Hybrid
  - One site determines the global schedule
  - Each site optimizes the local subqueries



# Network Topology

- **Wide area networks (WAN)** – point-to-point
  - Characteristics
    - Relatively low bandwidth (compared to local CPU/IO)
    - High protocol overhead
  - Communication cost may dominate; ignore all other cost factors
  - Global schedule to minimize communication cost
  - Local schedules according to centralized query optimization
- **Local area networks (LAN)**
  - Communication cost not that dominant
  - Total cost function should be considered
  - Broadcasting can be exploited (joins)
  - Special algorithms exist for star networks

# Distributed Query Processing Methodology



# Outline

- Distributed Query Processing
  - Query Decomposition and Localization
  - Distributed Query Optimization
  - Join Ordering
  - Adaptive Query Processing

# Step 1 – Query Decomposition

Same as centralized query processing

**Input :** Calculus query on global relations

- Normalization
  - Manipulate query quantifiers and qualification
- Analysis
  - Detect and reject “incorrect” queries
- Simplification
  - Eliminate redundant predicates
- Restructuring
  - Calculus query → algebraic query
  - Use transformation rules

## Step 2 – Data Localization

**Input:** Algebraic query on distributed relations

- Determine which fragments are involved
- **Localization program**
  - Substitute for each global query its materialization program
  - Optimize

# Example

- Assume
  - EMP is fragmented as follows:
    - $EMP_1 = \sigma_{ENO \leq "E3"}(EMP)$
    - $EMP_2 = \sigma_{"E3" < ENO \leq "E6"}(EMP)$
    - $EMP_3 = \sigma_{ENO \geq "E6"}(EMP)$
  - ASG fragmented as follows:
    - $ASG_1 = \sigma_{ENO \leq "E3"}(ASG)$
    - $ASG_2 = \sigma_{ENO > "E3"}(ASG)$
- In any query
  - Replace EMP by  $(EMP_1 \cup EMP_2 \cup EMP_3)$
  - Replace ASG by  $(ASG_1 \cup ASG_2)$

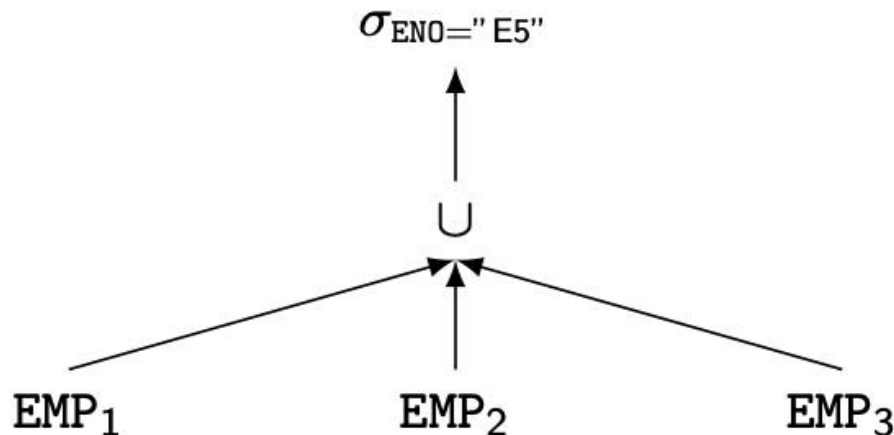
# Reduction for PHF

- Reduction with selection

- Relation  $R$  and  $F_R = \{R_1, R_2, \dots, R_w\}$  where  $R_j = \sigma_{p_j}(R)$

$$\sigma_{p_i}(R_j) = \emptyset \text{ if } \forall x \text{ in } R: \neg(p_i(x) \wedge p_j(x))$$

```
SELECT  *  
FROM    EMP  
WHERE   ENO="E5"
```



# Reduction for PHF

- Reduction with join
  - Possible if fragmentation is done on join attribute
  - Distribute join over union

$$(R_1 \cup R_2) \bowtie S \Leftrightarrow (R_1 \bowtie S) \cup (R_2 \bowtie S)$$

- Given  $R_i = \sigma_{p_i}(R)$  and  $R_j = \sigma_{p_j}(R)$

$$R_i \bowtie R_j = \emptyset \text{ if } \forall x \text{ in } R_i, \forall y \text{ in } R_j: \neg(p_i(x) \wedge p_j(y))$$

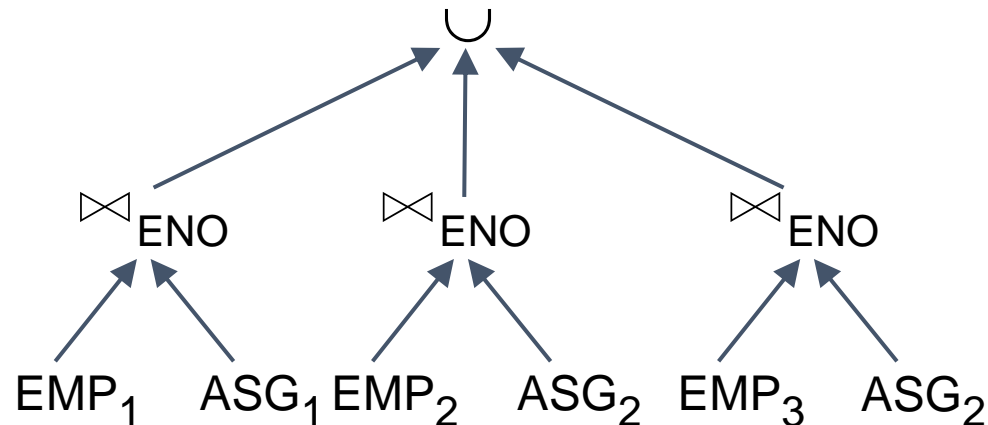



# Reduction for PHF

- Assume EMP is fragmented as before and
  - $ASG_1: \sigma_{ENO \leq "E3"}(ASG)$
  - $ASG_2: \sigma_{ENO > "E3"}(ASG)$
- Consider the query

**SELECT** \*  
**FROM** EMP  
**NATURAL JOIN** ASG

- Distribute join over unions
- Apply the reduction rule



# Reduction for VF

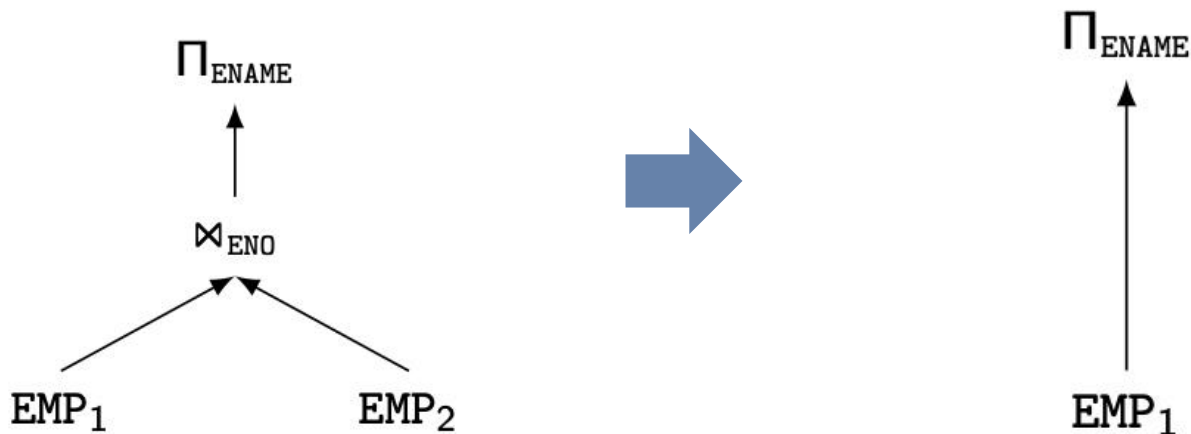
- Find useless (not empty) intermediate relations

Relation  $R$  defined over attributes  $A = \{A_1, \dots, A_n\}$  vertically fragmented as  $R_i = \Pi_{A'}(R)$  where  $A' \subseteq A$ :

$\Pi_{D,K}(R_i)$  is useless if the set of projection attributes  $D$  is not in  $A'$

Example:  $EMP_1 = \Pi_{ENO,ENAME}(EMP)$ ;  $EMP_2 = \Pi_{ENO,TITLE}(EMP)$

```
SELECT  ENAME
FROM    EMP
```



# Reduction for DHF

- Rule :
  - Distribute joins over unions
  - Apply the join reduction for horizontal fragmentation
- Example

$ASG_1: ASG \bowtie_{ENO} EMP_1$

$ASG_2: ASG \bowtie_{ENO} EMP_2$

$EMP_1: \sigma_{TITLE="Programmer"} (EMP)$

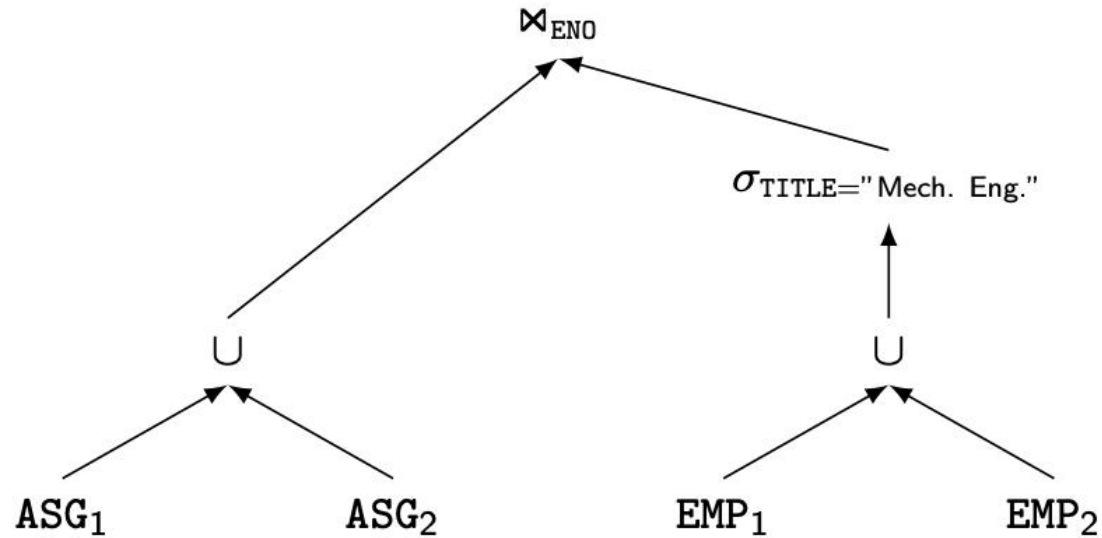
$EMP_2: \sigma_{TITLE!="Programmer"} (EMP)$

- Query


```
SELECT      *  
FROM  EMP NATURAL JOIN ASG  
WHERE EMP.TITLE = "Mech. Eng."
```

# Reduction for DHF

Generic query

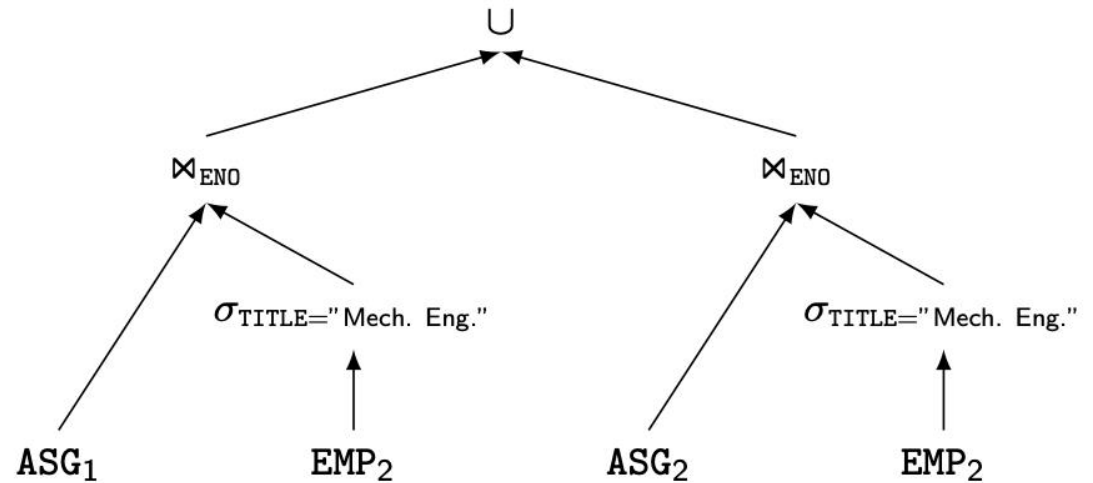


Selections first

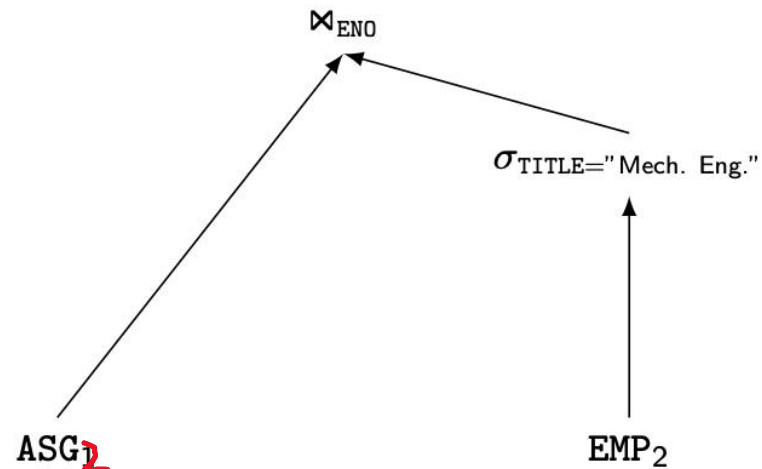


# Reduction for DHF

Joins over unions



Elimination of the empty  
intermediate relations  
(left sub-tree)



# Reduction for Hybrid Fragmentation

- Combine the rules already specified:
  - Remove **empty relations** generated by contradicting selections on horizontal fragments;
  - Remove **useless relations** generated by projections on vertical fragments;
  - Distribute **joins over unions** in order to isolate and remove useless joins.

# Reduction for HF

## Example

Consider the following hybrid fragmentation:

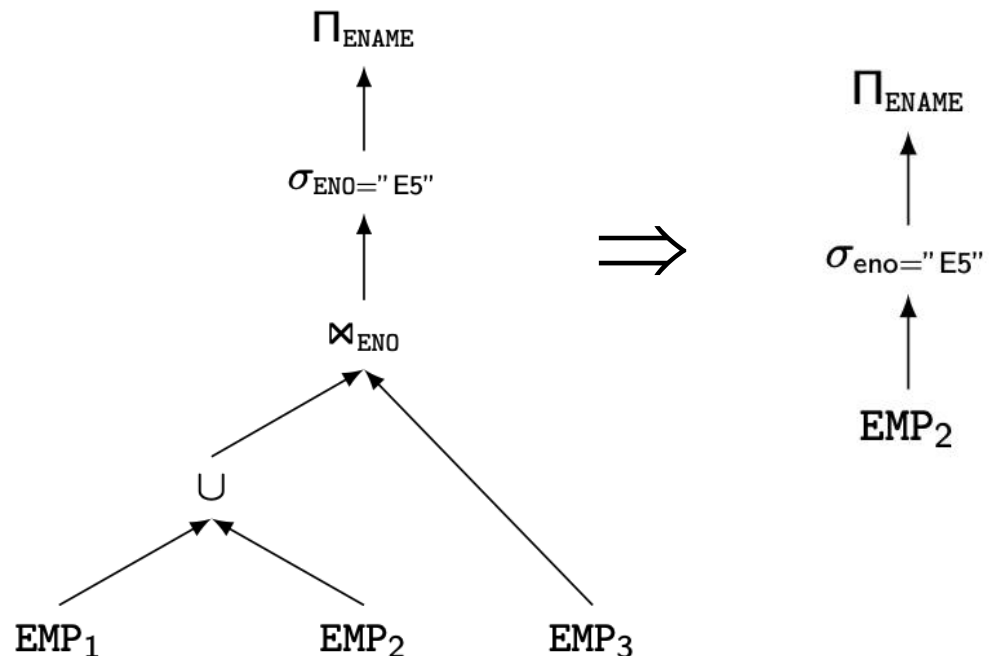
$$EMP_1 = \sigma_{ENO \leq "E4"} (\Pi_{ENO, ENAME} (EMP))$$

$$EMP_2 = \sigma_{ENO > "E4"} (\Pi_{ENO, ENAME} (EMP))$$

$$EMP_3 = \sigma_{ENO, TITLE} (EMP)$$

and the query

```
SELECT  ENAME
FROM    EMP
WHERE    ENO = "E5"
```



# Outline

- Distributed Query Processing
  - Query Decomposition and Localization
  - Distributed Query Optimization
  - Join Ordering
  - Adaptive Query Processing

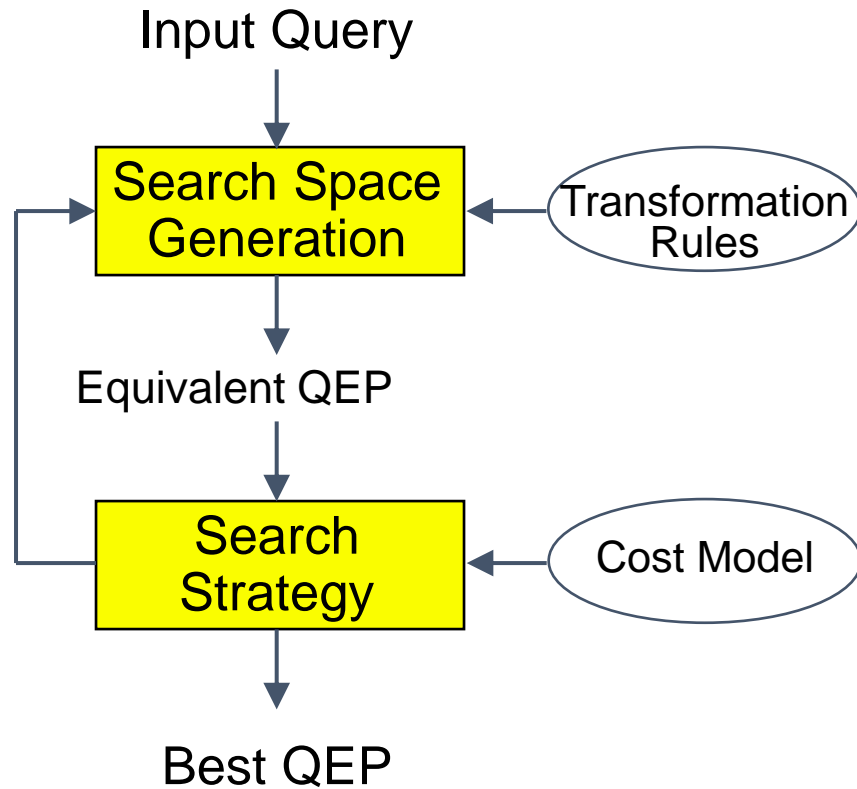


# Step 3 – Global Query Optimization

**Input:** Fragment query

- Find the *best* (not necessarily optimal) global schedule
  - Minimize a cost function
  - Distributed join processing
    - Bushy vs. linear trees
    - Which relation to ship where?
    - Ship-whole vs. ship-as-needed
  - Decide on the use of semi-joins
    - Semi-join saves on communication at the expense of more local processing
  - Join methods
    - Nested loop, merge join or hash join

# Query Optimization Process



# Components


- Search space
  - The set of equivalent algebra expressions (query trees)
- Cost model
  - I/O cost + CPU cost + communication cost
  - These might have different weights in different distributed environments (LAN vs WAN)
  - Can also maximize throughput
- Search algorithm
  - How do we move inside the solution space?
  - Exhaustive search, heuristic algorithms (iterative improvement, simulated annealing, genetic,...)

# Search Strategy

- How to “move” in the search space
- Deterministic
  - ➔ Start from base relations and build plans by adding one relation at each step
  - ➔ Dynamic programming: breadth-first
  - ➔ Greedy: depth-first
- Randomized
  - ➔ Search for optimalities around a particular starting point
  - ➔ Trade optimization time for execution time
  - ➔ Better when  $> 10$  relations
  - ➔ Simulated annealing
  - ➔ Iterative improvement

# Search Strategies

- Deterministic



- Randomized

